

ARIMA Modelling for a Univariate Time Series

July 15, 2024

0.0.1 Harsh Mittal

ARIMA Modelling

```
[28]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
[29]: df = pd.read_csv("C:/Users/harsh.hm.mittal/Downloads/MSFT.csv")
```

```
[30]: df.head()
```

```
[30]:
```

	Date	Close
0	14-07-2014	35.871395
1	15-07-2014	36.135273
2	16-07-2014	37.522812
3	17-07-2014	37.905876
4	18-07-2014	38.042068

```
[31]: df = df[["Close"]].copy()
```

```
[32]: df.describe()
```

```
[32]:
```

	Close
count	2516.000000
mean	164.301813
std	115.265938
min	34.823288
25%	57.289910
50%	129.766113
75%	254.843505
max	467.559998

0.0.2 Train test split

```
[33]: n = int(len(df) * 0.8)
      train = df.Close[:n]
      test = df.Close[n:]
```

```
[34]: print(len(train))
      print(len(test))
```

2012

504

0.0.3 ADF Test - Checking whether price series is stationary

```
[35]: from statsmodels.tsa.stattools import adfuller

      result = adfuller(train)
      print(f"ADF Statistics: {result[0]}")
      print(f"p-value: {result[1]}")
```

ADF Statistics: 0.17655902198335482

p-value: 0.9709157868214724


0.0.4 Autocorrelation Function (ACF)

```
[36]: from statsmodels.graphics.tsaplots import plot_acf
```

```
[37]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize = (16, 4))

      ax1.plot(train)
      ax1.set_title("Original")


      plot_acf(train, ax=ax2);
```



```
[38]: diff = train.diff().dropna()
```

```
fig, (ax1, ax2) = plt.subplots(1, 2, figsize = (16, 4))


ax1.plot(diff)
ax1.set_title("Difference once")
plot_acf(diff, ax=ax2);
```



```
[39]: diff = train.diff().diff().dropna()

fig, (ax1, ax2) = plt.subplots(1, 2, figsize = (16, 4))

ax1.plot(diff)
ax1.set_title("Difference twice")
plot_acf(diff, ax=ax2);
```



pmdarima package to get the number of differencing

```
[40]: !pip install pmdarima
```

```
Requirement already satisfied: pmdarima in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (2.0.4)
Requirement already satisfied: statsmodels>=0.13.2 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (0.13.2)
Requirement already satisfied: Cython!=0.29.18,!0.29.31,>=0.29 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (0.29.32)
```

```

Requirement already satisfied: numpy>=1.21.2 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (1.21.5)
Requirement already satisfied: scipy>=1.3.2 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (1.9.1)
Requirement already satisfied: scikit-learn>=0.22 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (1.0.2)
Requirement already satisfied: urllib3 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (1.26.11)
Requirement already satisfied: pandas>=0.19 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (1.4.4)
Requirement already satisfied: joblib>=0.11 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (1.1.0)
Requirement already satisfied: packaging>=17.1 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (21.3)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from pmdarima) (63.4.1)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from
packaging>=17.1->pmdarima) (3.0.9)
Requirement already satisfied: pytz>=2020.1 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from
pandas>=0.19->pmdarima) (2022.1)
Requirement already satisfied: python-dateutil>=2.8.1 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from
pandas>=0.19->pmdarima) (2.8.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from scikit-
learn>=0.22->pmdarima) (2.2.0)
Requirement already satisfied: patsy>=0.5.2 in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from
statsmodels>=0.13.2->pmdarima) (0.5.2)
Requirement already satisfied: six in
c:\users\harsh.hm.mittal\anaconda3\lib\site-packages (from
patsy>=0.5.2->statsmodels>=0.13.2->pmdarima) (1.16.0)

```

```
[41]: from pmdarima.arima.utils import ndiffs
```

```
[42]: ndiffs(train, test="adf")
```

```
[42]: 1
```

P

p is the order of the Auto Regressive (AR) term. It refers to the number of lags to be used as predictors.

Required number of AR terms can be found out by inspecting the Partial Autocorrelation (PACF) plot.


The partial autocorrelation represents the correlation between the series and its lags.

```
[43]: from statsmodels.graphics.tsaplots import plot_pacf
```

```
[44]: diff = train.diff().dropna()

fig, (ax1, ax2) = plt.subplots(1, 2, figsize = (16, 4))

ax1.plot(diff)
ax1.set_title("Difference once")
ax2.set_ylim(0, 1)
plot_pacf(diff, ax=ax2);
```



We can observe that the PACF lag 9 is significant as it's above the significance line.

q


q is the order of the Moving Average (MA) term. It refers to the number of lagged forecast errors that should go into the ARIMA model.

ACF plot is to be looked out for MA terms.

```
[45]: diff = train.diff().dropna()

fig, (ax1, ax2) = plt.subplots(1, 2, figsize = (16, 4))

ax1.plot(diff)
ax1.set_title("Difference once")
ax2.set_ylim(0, 1)
plot_acf(diff, ax=ax2);
```



0.0.5 Fitting the ARIMA model

```
[46]: from statsmodels.tsa.arima.model import ARIMA
```

```
# ARIMA Model
model = ARIMA(train, order=(9, 1, 9))
result = model.fit()
```

```
[47]: print(result.summary())
```

```

SARIMAX Results
=====
Dep. Variable:          Close    No. Observations:          2012
Model:                ARIMA(9, 1, 9)    Log Likelihood          -4802.858
Date:                Mon, 15 Jul 2024    AIC                    9643.716
Time:                20:19:26    BIC                    9750.237
Sample:                0    HQIC                    9682.817
                        - 2012
Covariance Type:          opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.1545	0.114	1.357	0.175	-0.069	0.378
ar.L2	0.6666	0.099	6.707	0.000	0.472	0.861
ar.L3	-0.3644	0.093	-3.923	0.000	-0.546	-0.182
ar.L4	-0.3521	0.111	-3.184	0.001	-0.569	-0.135
ar.L5	0.7325	0.059	12.472	0.000	0.617	0.848
ar.L6	-0.2822	0.092	-3.067	0.002	-0.463	-0.102
ar.L7	-0.6090	0.107	-5.684	0.000	-0.819	-0.399
ar.L8	0.2455	0.075	3.256	0.001	0.098	0.393
ar.L9	0.5351	0.077	6.927	0.000	0.384	0.686
ma.L1	-0.2764	0.116	-2.380	0.017	-0.504	-0.049
ma.L2	-0.6165	0.104	-5.911	0.000	-0.821	-0.412
ma.L3	0.4208	0.092	4.574	0.000	0.240	0.601
ma.L4	0.2749	0.112	2.446	0.014	0.055	0.495
ma.L5	-0.7440	0.054	-13.894	0.000	-0.849	-0.639
ma.L6	0.3122	0.090	3.481	0.000	0.136	0.488
ma.L7	0.6126	0.105	5.860	0.000	0.408	0.818
ma.L8	-0.3490	0.075	-4.636	0.000	-0.497	-0.201
ma.L9	-0.4270	0.078	-5.462	0.000	-0.580	-0.274
sigma2	6.9017	0.106	65.107	0.000	6.694	7.109

```

=====
Ljung-Box (L1) (Q):          0.67    Jarque-Bera (JB):
5424.09
```

```

Prob(Q):                                0.41    Prob(JB):
0.00
Heteroskedasticity (H):                  42.50    Skew:
-0.64
Prob(H) (two-sided):                    0.00    Kurtosis:
10.94
=====
===

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```

[48]: # Plot residual errors
residuals = pd.DataFrame(result.resid)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 4))


ax1.plot(residuals)
ax2.hist(residuals, density=True)

```

```

[48]: (array([2.72469915e-04, 2.54305254e-03, 1.31693792e-02, 1.51493273e-01,
              1.37143191e-02, 1.27152627e-03, 1.81646610e-04, 0.00000000e+00,
              0.00000000e+00, 9.08233051e-05]),
      array([-18.85221279, -13.37985201, -7.90749123, -2.43513045,
              3.03723033,  8.50959111, 13.98195188, 19.45431266,
              24.92667344, 30.39903422, 35.871395  ]),
      <BarContainer object of 10 artists>)

```



```

[49]: step = 30

forecast = result.get_forecast(steps=step)
fc = forecast.predicted_mean
se = forecast.se_mean
conf = forecast.conf_int()

```

```
[50]: fc = pd.Series(fc, index=test[:step].index)
lower = pd.Series(conf.iloc[:, 0], index=test[:step].index)
upper = pd.Series(conf.iloc[:, 1], index=test[:step].index)
```

```
[51]: plt.figure(figsize=(16, 8))
plt.plot(test[:step], label="actual")
plt.plot(fc, label="forecast")
plt.fill_between(lower.index, lower, upper, color="k", alpha=0.1)
plt.title("Forecast vs Actual")
plt.legend(loc="upper left")
```

```
[51]: <matplotlib.legend.Legend at 0x2195242a310>
```

